

Table of Contents for Programming Scala, 2nd Edition

 safaribooksonline.com/library/view/programming-scala-2nd/9781491950135/ch10.html

Chapter 10. The Scala Object System, Part I

We've learned a lot about Scala's implementation of object-oriented programming. In this chapter, we'll discuss more details of the type hierarchy in the standard library, exploring some of those types in depth, such as `Predef`.

But first, let's discuss an important feature of the type system called *variance under inheritance*, which we'll need to understand before discussing several of the library types described later in this chapter.

We'll conclude with a discussion of object equality.

Parameterized Types: Variance Under Inheritance

An important difference between Java's and Scala's parameterized types (usually called *generics* in the Java literature) is how *variance under inheritance* works.

For example, suppose a method takes an argument of type `List[AnyRef]`. Can you pass a `List[String]` value? In other words, should a `List[String]` be considered a *subtype* of `List[AnyRef]`? If true, this kind of variance is called *covariance*, because the supertype-subtype relationship of the container (the parameterized type) "goes in the same direction" as the relationship between the type parameters.

We can also have types that are *contravariant*, where `X[String]` is a *supertype* of `X[Any]`, for some type `X`.

If a parameterized type is neither covariant nor contravariant, it is called *invariant*. Conversely, some parameterized types can mix two or more of these behaviors.

Both Java and Scala support covariant, contravariant, and invariant types. However, in Scala, the variance behavior is defined as part of the type *declaration* using a so-called *variance annotation* on each type parameter, as appropriate. For covariant type parameters, `+` is used. For contravariant type parameters, `-` is used. No annotation is used for invariant type parameters. In other words, the type designer decides how the type should vary under inheritance.

Here are some example declarations (we'll see examples of real types shortly):

```
class W[+A] {...}           // covariant
class X[-A] {...}
// contravariant
class Y[A] {...}           // invariant
class Z[-A,B,+C] {...}     // mixed
```

In contrast, Java parameterized type *definitions* do not define the variance behavior under inheritance. Instead, the variance behavior of a parameterized type is specified when the type is *used*, i.e., at the *call site*, when variables are declared.

The three kinds of variance notations for Java and Scala and their meanings are summarized in [Table 10-1](#). `Tsup` is a *supertype* of `T` and `Tsub` is a *subtype* of `T`.

Table 10-1. Type variance annotations and their meanings

Scala	Java	Description
+T	? extends T	<i>Covariant</i> (e.g., <code>List[T_{sub}]</code> is a subtype of <code>List[T]</code>).
-T	? super T	<i>Contravariant</i> (e.g., <code>X[T^{sup}]</code> is a subtype of <code>X[T]</code>).
T	T	<i>Invariant</i> subclassing (e.g., can't substitute <code>Y[T^{sup}]</code> or <code>Y[T_{sub}]</code> for <code>Y[T]</code>).

Back to `List`—it is actually declared `List[+A]`, which means that `List[String]` is a subclass of `List[AnyRef]`, so `Lists` are covariant in the type parameter `A`. When a type like `List` has only one covariant type parameter, you'll often hear the shorthand expression “Lists are covariant” and similarly for contravariant types.

Covariant and invariant types are reasonably easy to understand. What about contravariant types?

Functions Under the Hood

The best example of contravariance is the set of traits `FunctionN`, such as `scala.Function2`, where `N` is between 0 and 22, inclusive, and corresponds to the number of arguments that a function takes. Scala uses these traits to implement anonymous functions.

We've been using anonymous functions, also known as function literals, throughout the book. For example:

```
List(1, 2, 3, 4) map (i => i + 3)
// Result: List(4, 5, 6, 7)
```

The function expression `i => i + 3` is actually *syntactic sugar* that the compiler converts to the following instantiation of an anonymous subclass of `scala.Function1`:

```
val f: Int => Int = new Function1[Int,Int] {
  def apply(i: Int): Int = i + 3
}
List(1, 2, 3, 4) map (f)
// Result: List(4, 5, 6, 7)
```

Note

The conventional name `apply` for the default method called when an object is followed by an argument list originated with the idea of *function application*. For example, once `f` is defined, we call it by applying an argument list to it, e.g., `f(1)`, which is actually `f.apply(1)`.

Historically, the JVM didn't allow “bare” functions in byte code. Everything had to be in an object wrapper. More recent versions of Java, especially Java 8, relax this restriction, but to enable Scala to work on older JVMs, the compiler has converted anonymous functions into anonymous subclasses of the appropriate `FunctionN` trait. You've probably written anonymous subclasses like this for Java interfaces many times in your Java projects.

The `FunctionN` traits are abstract, because the method `apply` is abstract. Note that we defined `apply` here. The

compiler does this for us when we use the more concise literal syntax instead, `i => i + 3`. That function body is used to define `apply`.

Note

Java 8 adds support for function literals, called *lambdas*. They use a different implementation than the one used by Scala, because Scala supports older JVMs.

Returning to contravariance, here is the declaration of `scala.Function2`:

```
trait Function2[-T1, -T2, +R] extends
AnyRef
```

The last type parameter, `+R`, is the return type. It is *covariant*. The leading two type parameters are for the first and second function arguments, respectively. They are *contravariant*. For the other `FunctionN` traits, the type parameters corresponding to function arguments are contravariant.

Therefore, functions have mixed variance behavior under inheritance.

What does this really mean? Let's look at an example to understand the variance behavior:

```
// src/main/scala/progscala2/objectsystem/variance/func.scX

class CSuper          { def msuper() = println("CSuper") }
// ❶

class C      extends CSuper { def m()      = println("C") }
class CSub   extends C      { def msub()   = println("CSub") }

var f: C => C = (c: C) => new C           // ❷
f          = (c: CSuper) => new CSub      // ❸
f          = (c: CSuper) => new C         // ❹
f          = (c: C)      => new CSub      // ❺
f          = (c: CSub)   => new CSuper    // ❻ COMPILATION ERROR!
```

❶

Define a three-level inheritance hierarchy of classes.

❷

We define one function `f` as a `var` so we can keep assigning new functions to it. All valid function instances

`C =>`
must be `C` (in other words, `Function1[C,C]`; note how we can use the literal syntax for the type, too).
The values we assign must satisfy the constraints of variance under inheritance for functions. The first
assignment is `(c: C) => new C` (it ignores the argument `c`).

❸

`(c: CSuper) => new CSub`, is valid, because the argument `C` is contravariant, so `CSuper` is a valid substitution, while the return value is covariant, so `CSub` is a valid replacement for `C`.

4

Similar to the previous case, but we simply return a `C`.

5

Similar to the previous cases, but we simply pass a `C`.

6

An error! A function that takes a `CSub` argument is invalid, because of contravariance, and the return type `CSuper` is invalid due to covariance.

This script doesn't produce any output. If you run it, it will fail to compile on the last line, but the other statements will be valid.

[Design by Contract](#) explains why these rules make sense. It is a formulation of the *Liskov Substitution Principle* and we'll discuss it briefly as a programming tool in [Better Design with Design by Contract](#). For now, let's try to understand intuitively why these rules work.

The function variable `f` is of type `C =>` (that is, `Function1[-C, +C]`). The first assignment to `f` matches this type signature exactly.

Now we assign different anonymous function values to `f`. The whitespace makes the similarities and differences stand out when comparing the original declaration of `f` and the subsequent reassignments. We keep reassigning to

`f` because we are testing what substitutions are valid for `C =>`.

The second assignment, `(x:CSuper) =>`, obeys the declaration for *contravariant* arguments and a *covariant* return type, but *why is this safe?*

The key insight is to recall how `f` will be used and what assumptions we can make about the actual function behind

it. When we say its type is `C =>`, we are defining a *contract* that any valid `C` value can be passed to `f` and `f` will never return anything other than a `C` value.

So, if the actual function has the type `(x:CSuper) =>`, that function not only accepts any `C` value as an argument, it can also handle any instance of the parent type `CSuper` or another of its other subtypes, if any. Therefore, because we only pass `C` instances, we'll never pass an argument to `f` that is outside the range of values it promises to accept. In a sense, `f` is more "permissive" than it needs to be for this use.

Similarly, when it returns only `CSub` instances, that is also safe, because the caller can handle instances of `C`, so it can certainly always handle instances of `CSub`. In this sense, `f` is more "restrictive" than it needs to be for this use.

The last line in the example breaks both rules for input and output types. Let's consider what would happen if we allowed that substitution for `f`.

In this case, the actual `f` would only know how to handle `CSub` instances, but the caller would believe that any `C` instance can be passed to `f`, so a runtime failure is likely when `f` is "surprised," i.e., it tries to call some method that is only defined for `CSub`, not `C`. Similarly, if the actual `f` can return a `CSuper`, it will "surprise" the caller with an instance that is outside the range of expected return values, the allowed instances of `C`.

This is why function arguments must be contravariant and return values must be covariant.

Variance annotations only make sense on the type parameters for *types*, not for parameterized methods, because the annotations affect the behavior of subtyping. Methods aren't subtyped. For example, the simplified signature for the `List.map` method looks like this:

```
sealed abstract class List[+A] ... {  
  // mixin traits  
  omitted  
  ...  
  def map[B] (f: A => B): List[B] = {...}  
  ...  
}
```

There is no variance annotation on `B` and if you tried to add one, the compiler would throw an error.

Note

The `+` *variance annotation* means the parameterized type is *covariant* in the type parameter. The `-` variance annotation means the parameterized type is *contravariant* in the type parameter. No variance annotation means the parameterized type is *invariant* in the type parameter.

Finally, the compiler checks your use of variance annotations for invalid uses. Here's what happens if you attempt to define your own function with the wrong annotations:

```
trait MyFunction2[+T1, +T2, -R]  
scala> {  
  def apply(v1:T1, v2:T2): R =  
    | ???  
    | }  
<console>:37: error: contravariant type R occurs in covariant position  
in type (v1: T1, v2: T2)R of method apply  
  def apply(v1:T1, v2:T2): R = ???  
    ^  
<console>:37: error: covariant type T1 occurs in contravariant position  
in type T1 of value v1  
  def apply(v1:T1, v2:T2): R = ???  
    ^  
<console>:37: error: covariant type T2 occurs in contravariant position  
in type T2 of value v2  
  def apply(v1:T1, v2:T2): R = ???  
    ^
```

Note the error messages. The compiler requires function arguments to behave *contravariantly* and return types to behave *covariantly*.

Variance of Mutable Types

All the parameterized types we've discussed so far have been immutable types. What about the variance behavior of mutable types? The short answer is that only *invariance* is allowed. Consider this example:

```
// src/main/scala/progscala2/objectsystem/variance/mutable-type-variance.scX
```

```
scala> class ContainerPlus[+A] (var value: A)
<console>:34: error: covariant type A occurs in contravariant position
in type A of value value_=
      class ContainerPlus[+A] (var value: A)
        ^

scala> class ContainerMinus[-A] (var value: A)
<console>:34: error: contravariant type A occurs in covariant position
in type => A of method value
      class ContainerMinus[-A] (var value: A)
        ^
```

The problem with a mutable field is that it behaves like a private field with public read *and* write accessor methods, even if the field is actually public and has no explicit accessor methods.

```
def value_=(newA: A): Unit = myinstance.value = someA
```

Recall from [Fields in Classes](#) that `Unit` + is the signature the compiler interprets as the setter for variable `value`. That is, we can write an expression `someA` and this method will be called. Note that the first error message uses this method signature and complains that we're using covariant type `A` in a contravariant position.

```
=>
```

The second error message mentions a method signature `A =>`. That is, a function that takes no arguments and returns an `A`, just like the *by-name* parameters we first saw in [Call by Name, Call by Value](#).

Here's another way to write the declaration using these methods explicitly, which looks more like traditional Java code:

```
class ContainerPlus[+A] (var a: A) {
  private var _value: A = a
  def value_=(newA: A): Unit = _value = newA
  def value: A = _value
}
```

```
value_=(newA: A)
```

Why must the `A` passed to `value_=(newA: A)` be contravariant? This doesn't seem right, because we're assigning a new value to `_value`, but if the new value can be a supertype of `A`, then we'll get a type error, because `_value` must be of type `A`, right?

Actually, that's the wrong way to think about the situation. The covariant/contravariant rules apply to how subclasses behave relative to superclasses.

Assume for a moment that our declaration is valid. For example, we could instantiate `ContainerPlus[C]`, using our `C`, `CSub`, and `CSuper` from before:

```

val cp = new ContainerPlus(new C)    //
❶
cp.value = new C                    //
❷
cp.value = new CSub                  //
❸
cp.value = new CSuper                //
❹

```

❶

Type parameter `A` is now `C`.

❷

Valid: we're just using the same type instance.

❸

Valid for the usual object-oriented reasons, since `CSub` is a subtype of `C`.

❹

Compilation error, because a `CSuper` instance can't be substituted for a `C` instance.

It's only when considering subtypes of `ContainerPlus` that trouble ensues:

```

val cp: ContainerPlus[C] = new ContainerPlus(new CSub)    //
❶
cp.value = new C                // ❷
cp.value = new CSub              // ❸
cp.value = new CSuper            // ❹

```

❶

Would be valid, if `ContainerPlus[+A]` were valid.

❷

From the declared type of `c`, this should be valid *and this is why the argument type must be contravariant*, but the actual `value_` method for the instance can't accept a `C` instance, because its `value` field is of type `CSub`.

❸

OK.

❹

OK.

The expression labeled ❷ illustrates why the method argument needs to be contravariant. The user of `c` expects the instance to work with `C` instances. By looking at the actual implementation of `value_`, we already know that we can't actually support contravariance, but let's ignore that for a moment and consider what happens if we change the variance annotation:

```
class ContainerMinus[-A] (var a: A) {
  private var _value: A = a
  def value_=(newA: A): Unit = _value = newA
  def value: A = _value
}
```

We already know from the error messages at the beginning of this section that this is considered OK for the `value_ =` method (even though it isn't actually OK), but now we get the second error we saw previously. The `A` is the return type for the `value` method, so `A` is in a covariant position.

Why must it be covariant? This is a little more intuitive. Again, the behavior of subtypes is the key. Once again, assume for a moment that the compiler allows us to instantiate `ContainerMinus` instances:

```
val cm: ContainerMinus[C] = new ContainerMinus(new CSuper) // ❶
val c: C = cm.value // ❷
val c: CSuper = cm.value // ❸
val c: CSub = cm.value // ❹
```

❶

Would be valid, if `ContainerMinus[-A]` were valid.

❷

`cm` thinks its `value` method returns a `C`, but the actual `value` method for the instance returns a `CSuper`.
Oops...

❸

OK.

❹

Fails for the same reason in line ❷.

So, if you think of a mutable field in terms of a getter and setter method, it appears in both covariant position when read and contravariant position when written. There is no such thing as a type parameter that is *both* contravariant and covariant, so the only option is for `A` to be invariant for the type of a *mutable* field.

Variance in Scala Versus Java

As we said, the variance behavior is defined in the *declaration* for Scala types, whereas it is defined when used, at the *call site*, in Java. The *client* of a type defines the variance behavior, defaulting to invariant. Java doesn't allow you to specify variance behavior at the definition site, although you can use expressions that look similar. Those expressions define *type bounds*, which we'll discuss shortly.

There are two drawbacks of Java's call-site variance specifications. First, it should be the library designer's job to understand the correct variance behavior and encode that behavior in the library itself. Instead, it's the library user who bears this burden. This leads to the second drawback. It's easy for a Java user to apply an incorrect annotation that results in unsafe code, like in the scenarios we just discussed.

Another problem in Java's type system is that `Arrays` are covariant in the type `T`. Consider this example:


```
// src/main/java/progscala2/objectsystem/JavaArrays.java
package progscala2.objectsystem;

public class JavaArrays {
    public static void main(String[] args) {
        Integer[] array1 = new Integer[] {
            new Integer(1), new Integer(2), new Integer(3) };
        // Compiles
        Number[] array2 = array1;      fine
        array2[2] = new Double(3.14);
        // Compiles, but throws a runtime
        error!
    }
}
```

This file compiles without error. However, when you run it with SBT, hilarity ensues:

```
> run-main progscala2.objectsystem.JavaArrays
[info] Running progscala2.objectsystem.JavaArrays
[error] (run-main-4) java.lang.ArrayStoreException: java.lang.Double
java.lang.ArrayStoreException: java.lang.Double
    at progscala2.objectsystem.JavaArrays.main(JavaArrays.java:10)
    ...
```

What's wrong? We discussed previously that *mutable* collections must be invariant in the type parameter to be safe. Because Java arrays are covariant, we're allowed by the compiler to assign an `Array[Integer]` instance to an `Array[Number]` reference. Then the compiler thinks it's OK to assign *any* `Number` to elements of the array, but in fact, the array "knows" internally that it can only accept `Integer` values (including subtype instances, if any), so it throws a runtime exception, defeating the purpose of static type checking. Note that even though Scala wraps Java `Arrays`, the Scala class `scala.Array` is invariant in the type parameter, so it prevents this "hole."

See Maurice Naftalin and Philip Wadler, *Java Generics and Collections*, O'Reilly Media, 2006 for more details of Java's generics and arrays, from which the last example was adapted.

The Scala Type Hierarchy

We already know many of the types in Scala's type hierarchy. Let's look at the general structure of the hierarchy and fill in more details. [Figure 10-1](#) shows the large-scale structure. Unless otherwise noted, all the types we'll discuss here are in the top-level `scala` package.

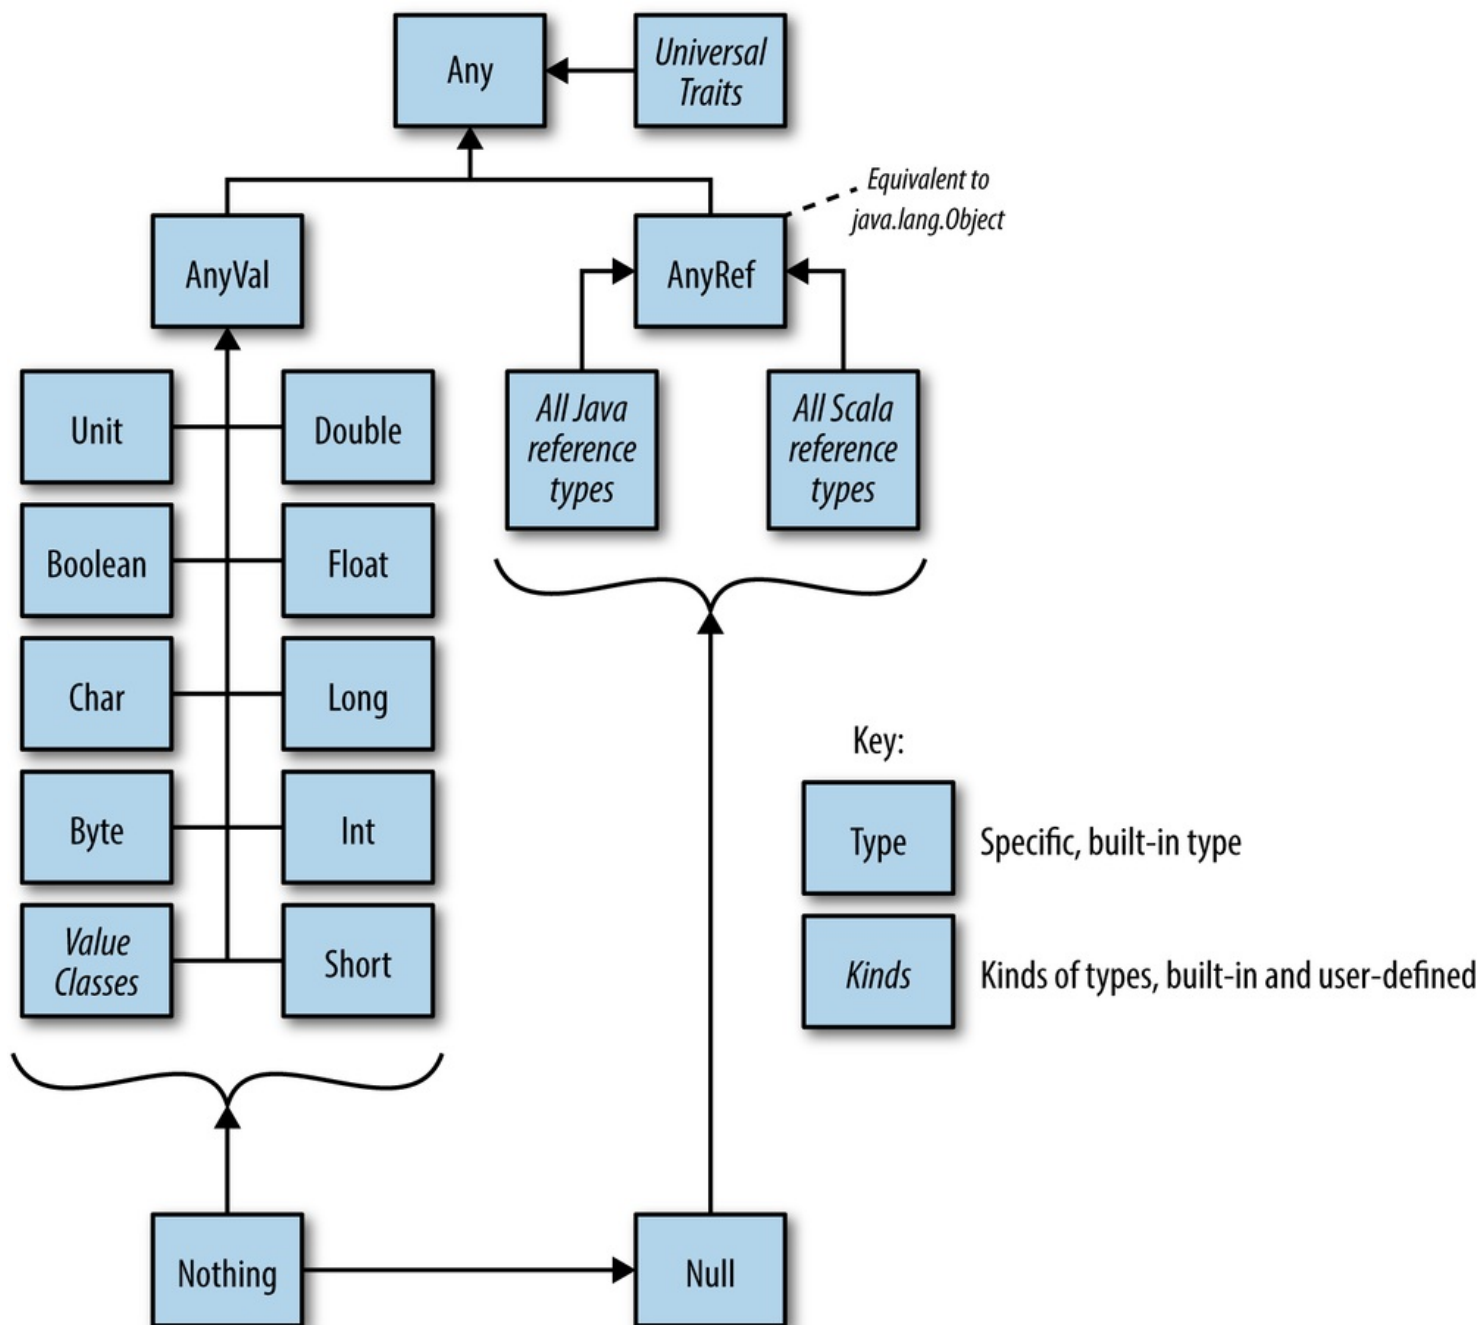


Figure 10-1. Scala's type hierarchy

At the root of the type hierarchy is **Any**. It has no parents and three children:

- **AnyVal**, the parent of value types and value classes
- **AnyRef**, the parent of all *reference* types
- *Universal Traits*, the newly introduced traits for the special uses we discussed in [Reference Versus Value Types](#)

AnyVal has nine concrete subtypes, called the *value types*. Seven are numeric value types: **Byte**, **Char**, **Short**, **Int**, **Long**, **Float**, and **Double**. The remaining two are nonnumeric value types, **Unit** and **Boolean**.

In addition, Scala 2.10 introduced user-defined *value classes*, which extend **AnyVal**, as discussed in [Reference Versus Value Types](#).

In contrast, all the other types are reference types. They are derived from `AnyRef`, which is the analog of `java.lang.Object`. Java's object model doesn't have a parent type of `Object` that encapsulates both primitives and reference types, because primitive types have special treatment.

Note

Before Scala 2.10, the compiler mixed into all instances of Scala reference types a “marker” trait named `ScalaObject`. The compiler no longer does this and the trait has been removed from Scala 2.11.

We have already learned about many of the reference types and we'll encounter more as we proceed. However, this is a good time to discuss a few of the widely used types.

Much Ado About Nothing (and Null)

`Nothing` and `Null` are two unusual types at the bottom of the type system. Specifically, `Nothing` is a subtype of all other types, while `Null` is a subtype of all reference types.

`Null` is the familiar concept from most programming languages, although they don't usually define a `Null` type, just a keyword `null` that's “assigned” to a reference to indicate the reference actually has no assigned value. `Null` is implemented in the compiler as if it has the following declaration:

```
package scala
abstract final class Null extends AnyRef
```

How can it be both `final` and `abstract`? This declaration disallows subtyping and creating your own instances, but the runtime environment provides one instance, the familiar `null` we know and love (cough, cough, ...).

`Null` is explicitly defined as a subtype of `AnyRef`, but it is also a subtype of all `AnyRef` types. This is the type system's formal way of allowing you to assign `null` to instances of any reference type. On the other hand, because `Null` is not a subtype of `AnyVal`, it is not possible to assign `null` to an `Int`, for example. Hence, Scala's `null` behaves exactly like Java's `null` behaves, as it must to coexist on the JVM. Otherwise, Scala could eliminate the concept of null and many potential bugs with it.

In contrast, `Nothing` has no analog in Java, but it fills a hole that exists in Java's type system. `Nothing` is implemented in the compiler as if it had the following declaration:

```
package scala
abstract final class Nothing extends Any
```

`Nothing` effectively extends `Any`. So, by construction in the type system, `Nothing` is a subclass of *all* other types, reference as well as value types. In other words, `Nothing` subclasses *everything*, as weird as that sounds.

Unlike `Null`, `Nothing` has no instances. Instead, it provides two capabilities in the type system that contribute to robust, type-safe design.

The first capability is best illustrated with our familiar `List[+A]` class. We now understand that `List` is *covariant* in `A`, so `List[String]` is a subtype of `List[Any]`, because `String` is a subtype of `Any`. Therefore, a `List[String]` instance can be assigned to a variable of type `List[Any]`.

Scala declares a type for the special case of an empty list, `Nil`. In Java, `Nil` would have to be a parameterized class like `List`, but this is unfortunate, because by definition, `Nil` never holds any elements, so a `Nil[String]`

and a `Nil[Any]` would be different, but without distinction.

Scala solves this problem by having `Nothing`. `Nil` is actually declared like this:

```
package scala.collection.immutable
object Nil extends List[Nothing] with Product with Serializable
```

We'll discuss `Product` in the next section. `Serializable` is the familiar Java “marker” interface for objects that can be serialized using Java's built-in mechanism for this purpose.

Note that `Nil` is an `object` and it extends `List[Nothing]`. There is only one instance of it needed, because it carries no “state” (elements). Because `List` is covariant in the type parameter, `Nil` is a subtype of `List[A]` for *all* types `A`. Therefore, we don't need separate `Nil[A]` instances. One will do.

`Nothing` and `Null` are called *bottom* types, because they reside at the bottom of the type hierarchy, so they are subtypes of all (or most) other types.

The other use for `Nothing` is to represent expressions that terminate the program, such as by throwing an exception. Recall the special `???` method in `Predef` we saw in [Nested Types](#). It can be called in a temporary method definition so the method is concrete, allowing an enclosing, concrete type to compile, but if the method is called, an exception is thrown. Here is the definition of `???`:

```
package scala
object Predef {
  ...
  def ??? : Nothing = throw new NotImplementedError
  ...
}
```

Because `???` “returns” `Nothing`, it can be called by any other function, no matter what type that function returns. Here is a pathological example:

```
scala> def m(l: List[Int]): List[Int] = l map (i => ???)
m: (l: List[Int])List[Int]

scala> m(List(1,2,3))
scala.NotImplementedError: an implementation is missing
  at scala.Predef$.qmark$qmark$qmark(Predef.scala:252)
  ...
```

Note that `m` is still expected to return a `List[Int]` and the definition type checks, even though `???` “returns” `Nothing`.

More realistically, `???` is called by a method that has been declared, but not yet defined:

```
/** @return (mean, standard_deviation)
 */
def mean_stdDev(data: Seq[Double]): (Double, Double) = ???
```

For normal termination, the `scala.sys` package defines an `exit` method, analogous to the `exit` method in Java's `System`. However, `sys.exit` returns `Nothing`.

This means that a method can declare that it returns a “normal” type, yet choose to call `sys.exit` if necessary, and still type check. A common example is the following idiom for processing command-line arguments, but exiting if an unrecognized option is provided:

```
// src/main/scala/progscala2/objectsystem/CommandArgs.scala
package progscala2.objectsystem

object CommandArgs {


  val help = """
  |usage: java ... objectsystem.CommandArgs
  |arguments
  |where the allowed arguments
  |are:
  | -h | --help          Show
  | help
  | -i | --in | --input path  Path for
  | input
  | -o | --on | --output path  Path for
  | input
  |""" .stripMargin

  ❶ def quit(message: String, status: Int): Nothing = { //
    if (message.length > 0) println(message)
    println(help)
    sys.exit(status)
  }

  ❷ case class Args(inputPath: String, outputPath: String) //

  def parseArgs(args: Array[String]): Args = {
    ❸ def pa(args2: List[String], result: Args): Args = args2 match { //
      case Nil => result
    // ❹
    case ("-h" | "--help") :: Nil => quit("", 0) //
    ❺
    case ("-i" | "--in" | "--input") :: path :: tail => //
    ❻
      pa(tail, result copy (inputPath = path))
    // ❼
    case ("-o" | "--out" | "--output") :: path :: tail => //
    ❽
      pa(tail, result copy (outputPath = path))
      "Unrecognized argument"
    case _ => quit(s"${args2.head}" //
    ❾
      , 1) //
  }

  val argz = pa(args.toList, Args("", ""))
}
```

```
// 
  if (argz.inputPath == "" || argz.outputPath == "") //
    "Must specify input and output
    quit(paths." , 1)
  argz
}

def main(args: Array[String]) = {
  val argz = parseArgs(args)
  println(argz)
}
}
```

1

Print an optional message, then the help message, then exit with the specified error status. Following Unix conventions, 0 is used for normal exits and nonzero values are used for abnormal termination. Note that `quit` returns `Nothing`.

2

A case class to hold the settings determined from the argument list.

3

A nested, recursively invoked function to process the argument list. We use the idiom of passing an `Args` instance to accumulate new settings (but by making a copy of it).

4

End of input, so return the accumulated settings.

5

User asks for help.

6

For the input argument, accept one of three variants for the option, `-i`, `--in`, or `--input`, followed by a path argument. Note that if the user doesn't provide a path (and there are no other arguments), the case won't match.

7

Call `pa` on the tail with an updated result.

8

Repeat for the output argument.

9

Handle the error of an unrecognized argument.

10

Call `pa` to process the arguments.

11

Verify that the input and output arguments were provided.

I find this example of pattern matching particularly elegant and concise.

This code is compiled when you build the project, so let's try it in `sbt`:

```
> run-main progscale2.objectsystem.CommandArgs -f
[info] Running progscale2.objectsystem.CommandArgs -f
Unrecognized argument -f

usage: java ... progscale2.objectsystem.CommandArgs arguments
where the allowed arguments are:
  -h | --help           Show help
  -i | --in | --input path Path for input
  -o | --on | --output path Path for input
```

```
Exception: sbt.TrapExitSecurityException thrown from the
  UncaughtExceptionHandler in thread "run-main-1"
  java.lang.RuntimeException: Nonzero exit code: 1
    at scala.sys.package$.error(package.scala:27)
[trace] Stack trace suppressed: run last compile:runMain for the full output.
[error] (compile:runMain) Nonzero exit code: 1
[error] ...
```

```
> run-main progscale2.objectsystem.CommandArgs -i foo -o bar
[info] Running progscale2.objectsystem.CommandArgs -i foo -o bar
Args(foo,bar)
[success] ...
```

We didn't throw an exception for the invalid argument, but `sbt` didn't like the fact that we called `exit`.

Products, Case Classes, and Tuples

Your case classes mix in the `scala.Product` trait, which provides a few generic methods for working with the fields of an instance, a `Person` instance, for example:

```
scala> case class Person(name: String, age: Int)
defined class Person

scala> val p: Product = Person("Dean", 29)
p: Product = Person(Dean,29)  //The case class instance is assignable to a Product
variable.

scala> p.productArity
res0: Int = 2                //The number of fields.

scala> p.productElement(0)
res1: Any = Dean            //Elements counted from zero.

scala> p.productElement(1)
res1: Any = 29

scala> p.productIterator foreach println
Dean
29
```

While having generic ways of accessing fields can be useful, its value is limited by the fact that `Any` is used for the fields' types, not their actual types.

There are also subtypes of `Product` for specific arities, up to 22 (for example, `scala.Product2` for two-element products). These types add methods for selecting a particular field, with the correct type information preserved. For example, `Product2[+T1, +T2]` adds these methods:

```
package scala
trait Product2[+T1, +T2] extends Product
{
  abstract def _1: T1
  abstract def _2: T2
  ...
}
```

These methods return the actual type of the field. The type parameters are covariant, because the `ProductN` traits are only used with immutable types, where reading the field with a method like `_1` uses the corresponding type parameter, `T1`, in covariant position (i.e., only as a return type).

Recall that these methods are the same ones used to access the elements of tuples. In fact, all `TupleN` types extend the corresponding `ProductN` trait and provide concrete implementations for these `_1` to `_N` methods, for `_N` up to and including 22:


```
scala> val t2 = ("Dean", 29)
t2: (String, Int) = (Dean,29)

scala> t2._1
res0: String = Dean

scala> t2._2
res2: Int = 29

scala> t2._3
<console>:36: error: value _3 is not a member of (String, Int)
      t2._3
          ^
```

There is no third element nor a `_3` method for `Tuple2`.

Why the upper limit of 22? It's really somewhat arbitrary and you could make the reasonable argument that having 22 elements in a tuple is way too many anyway.

That's true for human comprehension, but unfortunately, there is a common scenario where exceeding this limit would be useful: in holding the fields (or columns) of large data "records." It's not uncommon for SQL and NoSQL data sets to have schemas with more than 22 elements. Tuples are useful, at least for smaller schemas, because they preserve the field (column) order and type. Hence, a 22-element limit is a problem.

It turns out that case classes in Scala 2.10 are also limited to 22 fields, but this implementation restriction has been eliminated in 2.11. Hence, data applications can use a case class for records with more than 22 elements.

Note

In Scala 2.10 and earlier, case classes are limited to 22 fields or less. This limitation is removed in Scala 2.11.

Hopefully, the 22-element limit for traits and products will be removed in a future release of Scala.

The Predef Object

For your convenience, whenever you compile code, the Scala compiler automatically imports the definitions in the top-level Scala package, called `scala`, as well as the definitions in the `java.lang` package (just like `javac`). Hence, many common Java and Scala types can be used without explicitly importing them or using fully qualified names. In addition, the compiler imports the definitions in the `Predef` object that provides a number of useful definitions, many of which we've discussed previously.

Let's round out the features provided by `Predef`. Note that a number of changes are being introduced in the Scala 2.11 version of `Predef`, most of which are not visible. We'll discuss the 2.11 version.

Implicit Conversions

First, `Predef` defines many implicit conversions. One group of conversions includes wrappers around the `AnyVal` types:

```

@inline implicit def byteWrapper(x: Byte)      = new runtime.RichByte(x)
@inline implicit def shortWrapper(x: Short)    = new runtime.RichShort(x)
@inline implicit def intWrapper(x: Int)        = new runtime.RichInt(x)
@inline implicit def charWrapper(c: Char)      = new runtime.RichChar(c)
@inline implicit def longWrapper(x: Long)      = new runtime.RichLong(x)
@inline implicit def floatWrapper(x: Float)    = new runtime.RichFloat(x)
@inline implicit def doubleWrapper(x: Double)  = new runtime.RichDouble(x)
@inline implicit def booleanWrapper(x: Boolean) = new runtime.RichBoolean(x)

```

The `Rich*` types add additional methods, like comparison methods such as `<=` and `compare`. The `@inline` annotation encourages the compiler to “inline” the method call, i.e., eliminate it by inserting the

```

new
runtime.RichY(x)      logic directly.

```

Why have two separate types for bytes, for example? Why not put all the methods in `Byte` itself? The reason is that the extra methods would force an instance to be allocated on the heap, due to implementation requirements for byte code. `Byte` instances, like the other `AnyVal` types, are not actually heap-allocated, but represented as Java’s `byte` primitives. So, having separate `Rich*` types avoids the heap allocation except for those times when the extra methods are needed.

There are also methods for wrapping Java’s mutable arrays in instances of

`scala.collection.mutable.WrappedArray`, which adds many of the collection methods we discussed in [Chapter 6](#):

```

implicit def wrapIntArray(xs: Array[Int]): WrappedArray[Int]
implicit def wrapDoubleArray(xs: Array[Double]): WrappedArray[Double]
implicit def wrapLongArray(xs: Array[Long]): WrappedArray[Long]
implicit def wrapFloatArray(xs: Array[Float]): WrappedArray[Float]
implicit def wrapCharArray(xs: Array[Char]): WrappedArray[Char]
implicit def wrapByteArray(xs: Array[Byte]): WrappedArray[Byte]
implicit def wrapShortArray(xs: Array[Short]): WrappedArray[Short]
implicit def wrapBooleanArray(xs: Array[Boolean]): WrappedArray[Boolean]
implicit def wrapUnitArray(xs: Array[Unit]): WrappedArray[Unit]

```

Why are there separate methods for each `AnyVal` type? Each one uses a custom subclass of `WrappedArray` that exploits the fact that Java arrays of primitives are more efficient than arrays of boxed elements, so the less efficient, generic implementation for reference types is avoided.

There are similar methods for converting to `scala.collection.mutable.ArrayOps`. The only difference between `WrappedArray` and `ArrayOps` is that transformation functions for `WrappedArray`, such as `filter`, will return a new `WrappedArray`, while the corresponding functions in `ArrayOps` return `Arrays`.

Like `WrappedArray` and `ArrayOps`, there are analogous types for `Strings`, `scala/collection/immutable/WrappedString` and `scala/collection/immutable/StringOps`, which add collection methods to `Strings`, treating them like collections of `Chars`. Hence, `Predef` defines conversions between `String` and these types:

```
implicit def wrapString(s: String): WrappedString
implicit def unwrapString(ws: WrappedString): String

implicit def augmentString(x: String): StringOps
implicit def unaugmentString(x: StringOps): String
```

Note

Having pairs of similar wrapper types, like `WrappedArray/ArrayOps` and `WrappedString/StringOps`, is a bit confusing, but fortunately the implicit conversions are invoked automatically, selecting the correct wrapper type for the method you need.

There are other methods for converting between Java's boxed types for primitives and Scala's `AnyVal` types. They make Java interoperability easier:

```
implicit def byte2Byte(x: Byte)           = java.lang.Byte.valueOf(x)
implicit def short2Short(x: Short)        = java.lang.Short.valueOf(x)
implicit def char2Character(x: Char)       = java.lang.Character.valueOf(x)
implicit def int2Integer(x: Int)           = java.lang.Integer.valueOf(x)
implicit def long2Long(x: Long)           = java.lang.Long.valueOf(x)
implicit def float2Float(x: Float)        = java.lang.Float.valueOf(x)
implicit def double2Double(x: Double)     = java.lang.Double.valueOf(x)
implicit def boolean2Boolean(x: Boolean)   = java.lang.Boolean.valueOf(x)

implicit def Byte2byte(x: java.lang.Byte): Byte           = x.byteValue
implicit def Short2short(x: java.lang.Short): Short       = x.shortValue
implicit def Character2char(x: java.lang.Character): Char  = x.charValue
implicit def Integer2int(x: java.lang.Integer): Int        = x.intValue
implicit def Long2long(x: java.lang.Long): Long            = x.longValue
implicit def Float2float(x: java.lang.Float): Float        = x.floatValue
implicit def Double2double(x: java.lang.Double): Double    = x.doubleValue
implicit def Boolean2boolean(x: java.lang.Boolean): Boolean = x.booleanValue
```

Finally, in Scala 2.10, there is a group of implicit conversions that *prevent* `null` from being accepted as a value for an assignment. We'll just show one example, for `Byte`:

```
implicit def Byte2byteNullConflict(x: Null): Byte = sys.error(error"          "value
                                     "          ")
```

It triggers the following error:

```
scala> val b: Byte = null
<console>:23: error: type mismatch;
  found   : Null(null)
  required: Byte
Note that implicit conversions are not applicable because they are ambiguous:
  both method Byte2byteNullConflict in class LowPriorityImplicits of
  type (x: Null)Byte and method Byte2byte in object Predef of type (x: Byte)Byte
  are possible conversion functions from Null(null) to Byte
    val b: Byte = null
                ^
```

That works OK, but the error message isn't really that clear. It's complaining about ambiguous implicits, which the library deliberately introduced, but really it should just tell us more directly that we shouldn't assign `null` to anything.

Here's the error message produced by Scala 2.11:

```
scala> val b: Byte = null
<console>:7: error: an expression of type Null is ineligible for
  implicit conversion
    val b: Byte = null
                ^
```

Scala 2.11 eliminates the conversion methods and provides a better, more concise error message.

Type Definitions

`Predef` defines several types and type aliases.

To encourage the use of immutable collections, `Predef` defines aliases for the most popular, immutable collection types:

```
type Map[A, +B]      = collection.immutable.Map[A, B]
type Set[A]          = collection.immutable.Set[A]
type Function[-A, +B] = Function1[A, B]
```

Two convenient aliases for two- and three-element tuples have been deprecated in 2.11, on the grounds that they aren't used enough and don't add enough value to justify their existence:

```
type Pair[+A, +B]    = Tuple2[A, B]
type Triple[+A, +B, +C] = Tuple3[A, B, C]
]
```

Other `Predef` type members support type inference:

```
final class ArrowAssoc[A] extends
  AnyVal
```

```
    a ->
```

Used to implement the `b` literal syntax for creating two-element tuples. We discussed it in [Implicit](#)

Conversions.

```
sealed abstract class <:[-From, +To] extends (From) => To with
Serializable
```

Witnesses that type `From` is a subtype of type `To`. We discussed it in [Implicit Evidence](#).

```
sealed abstract class ==[-From, +To] extends (From) => To with
Serializable
```

Witnesses that types `From` and `To` are equal. We also mentioned it in [Implicit Evidence](#).

```
type Manifest[T] =
reflect.Manifest[T]
```

Used to retain type information that's lost in the JVM's *type erasure*. There is a similar type `OptManifest`.

We'll discuss them in [Class Tags, Type Tags, and Manifests](#).

Other types, like `scala.collection.immutable.List`, are made visible through nested imports within `Predef`. Also, companion objects for some of the types are also made visible, such as `==`, `Map`, and `Set`.

Condition Checking Methods

Sometimes you want to assert a condition is true, perhaps to “fail fast” and especially during testing. `Predef` defines a number of methods that assist in this goal:

```
def assert(assertion:
Boolean)
```

Test that `assertion` is true. If not, throw a `java.lang.AssertionError`.

```
def assert(assertion: Boolean, message: =>
Any)
```

Similar, but with an additional argument that's converted to a string message.

```
def assume(assertion:
Boolean)
```

Identical to `assert`, but conveys the meaning that the condition is assumed to be true when entering a block of code, such as a method.

```
def assume(assertion: Boolean, message: =>
Any)
```

Similar, but with an additional argument that's converted to a string message.

```
def require(requirement:
Boolean)
```

Identical to `assume`, but the Scaladoc says it conveys the meaning that the caller failed to satisfy some requirement; it could also convey the meaning that an implementation could not achieve a required result.

```
def require(requirement: Boolean, message: =>
Any)
```

Similar, but with an additional argument that's converted to a string message.

It's not shown, but all of these assertion methods are annotated with `@elidable(ASSERTION)`. The `@elidable` annotation tells the compiler not to generate byte code for a definition unless the argument to the annotation (`ASSERTION` in this case) is above a threshold that is specified during compilation. For example,

```
scalac -Xelide-below
2000
```

suppresses code generation for all annotated definitions with argument values below 2000. 2000 happens to be the value defined for `ASSERTION` in the `elidable` companion object. See the [Scaladoc page](#) for more information on `@elidable`.

Input and Output Methods

We've enjoyed the convenience of writing `println("foo")` instead of the more verbose Java equivalent, `System.out.println("foo")`. `Predef` gives us four variants for writing strings to `stdout`:

```
def print(x: Any):
Unit
    Convert x to a String and write it to stdout, without adding a line feed at the end automatically.
def printf(format: String, xs: Any*):
Unit
    Format a printf-style string using format as the format and the rest of the arguments xs, then write the
    resulting String to stdout, without adding a line feed at the end automatically.
def println(x: Any):
Unit
    Like print, but appends a line feed at the end automatically.
def println():
Unit
    Writes a blank line to stdout.
```

`Predef` in Scala 2.10 also defines several functions for reading input from `stdin`. However, these functions are deprecated in Scala 2.11. Instead, they are defined in a new `scala.io.ReadStdin` object that should be used instead. Otherwise, the method signatures and behaviors are the same:

```
def readBoolean():
Boolean
    Reads a Boolean value from an entire line from stdin.
def readByte():
Byte
    Reads a Byte value from an entire line from stdin.
def readChar():
Char
    Reads a Char value from an entire line from stdin.
def readDouble():
Double
    Reads a Double value from an entire line from stdin.
def readFloat():
Float
    Reads a Float value from an entire line from stdin.
def readInt():
Int
    Reads an Int value from an entire line from stdin.
def readLine(text: String, args: Any*):
String
    Prints formatted text to stdout and reads a full line from stdin.
def readLine():
String
    Reads a full line from stdin.
def readLong():
Long
    Reads a Long value from an entire line from stdin.
def readShort():
Short
    Reads a Short value from an entire line from stdin.
def readf(format: String):
List[Any]
    Reads in structured input from stdin as specified by the format specifier.
def readf1(format: String):
Any
    Reads in structured input from stdin as specified by the format specifier, returning only the first value
```

extracted, according to the format specification.

```
def readf2(format: String): (Any, Any)
```

Reads in structured input from `stdin` as specified by the `format` specifier, returning only the first two values extracted, according to the format specification.

```
def readf3(format: String): (Any, Any, Any)
```

Reads in structured input from `stdin` as specified by the `format` specifier, returning only the first three values extracted, according to the format specification.

Miscellaneous Methods

Finally, there are a few more useful methods in `Predef` to highlight:

```
def ???:  
Nothing
```

Called in a method body for a method that is actually unimplemented. It provides a concrete definition for the method, allowing enclosing types to compile as concrete (as opposed to abstract). However, if called, the method throws a `scala.NotImplementedError`. We first discussed it in [Nested Types](#).

```
def identity[A](x: A):  
A
```

Simply returns the argument `x`. It is useful for passing to *combinator* methods when no change is required. For example, a work flow calls `map` to transform the elements of a collection, passing it a configurable function to do the transformation. Sometimes, no transformation is required, so you'll pass `identity` instead.

```
def implicitly[T](implicit e: T):  
T
```

Used when an implicit argument list is specified with the type shorthand `M` ^[T : M], in which case the compiler adds an implicit argument list of the form `(implicit arg: M[T])`. (The actual name isn't `arg`, but something unique synthesized by the compiler.) Calling `implicitly` returns the argument `arg`. Discussed previously in [Using implicitly](#).

Now let's consider a very important topic in object-oriented design, checking the equality of objects.

Equality of Objects

Implementing a reliable equality test for instances is difficult to do correctly. Joshua Block's popular book, *Effective Java* (Addison-Wesley), and the *Scaladoc* page for [AnyRef.eq](#) describe the requirements for a good equality test.

Martin Odersky, Lex Spoon, and Bill Venners wrote a very good article on writing `equals` and `hashCode` methods, [How to Write an Equality Method in Java](#). Recall that these methods are created automatically for case classes.

In fact, I *never* write my own `equals` and `hashCode` methods. I find that any object that I use where I might need to test for equality or to use it as a `Map` key (where `hashCode` is used) *should* be a `case` class!

Caution

Some of the equality methods have the same names as equality methods in other languages, but the semantics are sometimes different!

Let's look at the different methods used to test equality.

The equals Method

We'll use a case class to demonstrate how the different equality methods work:

```
// src/main/scala/progscala2/objectsystem/person-equality.sc

case class Person(firstName: String, lastName: String, age: Int)

val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2   = Person("Buck", "Trends", 30)
```

The `equals` method tests for *value* equality. That is, `obj1 equals obj2` is true if both `obj1` and `obj2` have the same value. They do not need to refer to the same instance:

```

// =
p1a equals p1a      true
// =
p1a equals p1b      true
// =
p1a equals p2       false
// =
p1a equals null     false
null equals p1a
// throws
java.lang.NullPointerException
null equals null
// throws
java.lang.NullPointerException
```

Hence, `equals` behaves like the `equals` method in Java and the `eq?` method in Ruby, for example.

The == and != Methods

Whereas `==` is an operator in many languages, it is a method in Scala. In Scala 2.10, `==` is defined as `final` in `Any` and it delegates to `equals`. A different implementation is used in 2.11, but the behavior is effectively the same:

```

// =
p1a == p1a      true
// =
p1a == p1b      true
p1a == p2
// =
false
```

Hence, it behaves exactly like `equals`, namely it tests for *value* equality. The exception is the behavior when `null` is on the lefthand side:


```

// =
pla == null      false
// =
null == pla      false
null == null
// = true (compiler warns that it's always
true)

```

Should `null == null` be true? Actually, a warning is emitted:

```

<console>:8: warning: comparing values of types Null and Null using
`=='
will always yield true

```

As you would expect, `!=` is the negation, i.e., it is equivalent to `!(obj1 == obj2)`:

```

// =
pla != pla      false
// =
pla != plb      false
// =
pla != p2       true
// =
pla != null     true
// =
null != pla     true
null != null
// = false (compiler warns that it's always
false.)

```

Note

In Java, C++, and C#, the `==` operator tests for *reference*, not *value* equality. In contrast, Scala's `==` operator tests for *value* equality.

The eq and ne Methods

The `eq` method tests for *reference* equality. That is, `obj1 eq obj2` is true if both `obj1` and `obj2` point to the same location in memory. These methods are only defined for `AnyRef`:

```

pla eq pla      // = true
pla eq plb      // = false
pla eq p2       // = false
pla eq null     // = false
null eq pla     // = false
null eq null    // = true (compiler warns that it's always
true.)

```

Just as the compiler issues a warning for `null == null eq`, it issues the same warning for `null eq null`.

Hence, `eq` behaves like the `==` operator in Java, C++, and C#.

The `ne` method is the negation of `eq`, i.e., it is equivalent to `!(obj1 eq obj2)`:

```
// =
pla ne pla      false
// =
pla ne plb      true
// =
pla ne p2       true
// =
pla ne null     true
// =
null ne pla     true
null ne null
// = false (compiler warns that it's always
false.)
```

Array Equality and the `sameElements` Method

Comparing the contents of two `Arrays` doesn't have an obvious result in Scala:

```
Array(1, 2) == Array(1, 2)
// =
false
```

That's a surprise! Thankfully, there's a simple solution in the form of the `sameElements` method:

```
Array(1, 2) sameElements Array(1, 2)
// =
true
```

Actually, it's better to remember that `Arrays` are the mutable, raw Java arrays we know and love, which don't have the same methods we're accustomed to in the Scala collections library.

So, if you're tempted to compare arrays, consider whether or not it would be better to work with sequences instead (an argument for *not* using alternatives is when you really need the performance benefits of arrays).

In contrast, for example, `Lists` work as you would expect:

```
List(1, 2) == List(1, 2)
// =
true
List(1, 2) sameElements List(1, 2)
// =
true
```

Recap and What's Next

We discussed important topics in Scala's object system, such as the behavior under inheritance, the features in `Predef`, the fundamentals of the type hierarchy, and equality.

Next we'll complete our discussion of the object system by examining the behavior of member overriding and resolution rules.